



SCEDULING POLICY: REVIEW

Workload

JOB	arrival	run
A	0	40
B	0	20
C	5	10

Timelines





Schedulers:

FIFO

SJF

STCF

RR



UNIVERSITY of WISCONSIN-MADISON
Computer Sciences Department

CS 537
Introduction to Operating Systems

Andrea C. Arpaci-Dusseau
Remzi H. Arpaci-Dusseau

MEMORY VIRTUALIZATION:

Questions answered in this lecture:

What is in the address space of a process (review)?

What are the different ways that that OS can virtualize memory?

Time sharing, static relocation, dynamic relocation

(base, base + bounds, segmentation)

What hardware support is needed for dynamic relocation?

ANNOUNCEMENTS

- Switching discussion sections unofficially
- Project 1
 - Handin directories now available
 - Some test scripts available
 - Due next Friday, 5pm
- Lots of Office Hours

MORE VIRTUALIZATION

1st part of course: Virtualization

Virtual CPU: *illusion of private CPU registers*


- 2 lectures (mechanism + policy)

Virtual RAM: *illusion of private memory*

- 5 lectures

MOTIVATION FOR VIRTUALIZING MEMORY

Uniprogramming: One process runs at a time



Disadvantages:

- Only one process runs at a time
- Process can destroy OS

MULTIPROGRAMMING GOALS FOR MEMORY

Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

Efficiency

- Do not waste memory resources (minimize fragmentation)

Sharing

- Cooperating processes can share portions of address space

ABSTRACTION: ADDRESS SPACE

Address space: Each process has set of addresses


Problem:

How can OS provide illusion of private address space to each process?

Review: What is in an address space?

Address space has static and dynamic components

- Static: Code and some global variables
- Dynamic: Stack and Heap



MOTIVATION FOR DYNAMIC MEMORY

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time
- Must be pessimistic when allocate memory statically
 - Allocate enough for worst possible case; Storage is used inefficiently

Recursive procedures

- Do not know how many times procedure will be nested

Complex data structures: lists and trees

- `struct my_t *p = (struct my_t *)malloc(sizeof(struct my_t));`

Two types of dynamic allocation

- Stack
- Heap

STACK ORGANIZATION

Definition: Memory must be freed in opposite order from allocation

```
alloc(A);
alloc(B);
alloc(C);
free(C);
alloc(D);
free(D);
free(B);
free(A);
```

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

WHERE ARE STACKS USED?

OS uses stack for procedure call frames (local variables and parameters)

```
main () {
    int A = 0;
    foo (A);
    printf("A: %d\n", A);
}

void foo (int Z) {
    int A = 2;
    Z = 5;
    printf("A: %d Z: %d\n", A, Z);
}
```

HEAP ORGANIZATION

Definition: Allocate from any random location: malloc(), new()


- Heap memory consists of allocated areas and free areas (holes)
- Order of allocation and free is unpredictable

Advantage

- Works for all data structures

Disadvantages

- Allocation can be slow
- End up with small chunks of free space - fragmentation
- Where to allocate 12 bytes? 16 bytes? 24 bytes??
- What is OS's role in managing heap?
- OS gives big chunk of free memory to process; library manages individual allocations



QUIZ: MATCH THAT ADDRESS LOCATION

```
int x;
int main(int argc, char *argv[]) {
    int y;
    int *z = malloc(sizeof(int));
}
```

Possible segments: static data, code, stack, heap

What if no static data segment?

Address	Location
x	Static data → Code
main	Code
y	Stack
z	Stack
*z	Heap

MEMORY ACCESSES

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int x;
    x = x + 3;
}
```

`otool -tv demo1.o
(or objdump on Linux)`



```
0x10: movl 0x8(%rbp), %edi
0x13: addl $0x3, %edi
0x19: movl %edi, 0x8(%rbp)
```

%rbp is the base pointer:
points to base of current stack frame

QUIZ: MEMORY ACCESSES?

Initial **%rip** = 0x10
%rbp = 0x200

→ 0x10: `movl 0x8(%rbp), %edi`
 0x13: `addl $0x3, %edi`
 0x19: `movl %edi, 0x8(%rbp)`

%rbp is the base pointer:
points to base of current stack frame

%rip is instruction pointer (or program counter, PC)

1) Fetch instruction at addr 0x10
Exec:

2) load from addr 0x208

3) Fetch instruction at addr 0x13
Exec:

no memory access

4) Fetch instruction at addr 0x19
Exec:

5) store to addr 0x208

Memory Accesses to what addresses?

HOW TO VIRTUALIZE MEMORY?

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions across different processes?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation



1) TIME SHARING OF MEMORY

Try similar approach to how OS virtualizes CPU



Observation:



OS gives illusion of many virtual CPUs by saving **CPU registers to memory** when a process isn't running



Could give illusion of many virtual memories by saving **memory to disk** when process isn't running



TIME SHARE MEMORY: EXAMPLE







PROBLEMS WITH TIME SHARING MEMORY

Problem: Ridiculously poor performance

Better Alternative: **space sharing**

- At same time, space of memory is divided across processes

Remainder of solutions all use space sharing

REPEAT: HOW TO VIRTUALIZE MEMORY?

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions across different processes?

Possible Solutions for Mechanisms (covered today):


1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation

2) STATIC RELOCATION

Idea: OS rewrites each program before loading it as a process in memory

Each rewrite for different process uses different addresses and pointers

Change jumps, loads of static data



HOW TO VIRTUALIZE MEMORY?

Problem: How to run multiple processes simultaneously?


Addresses are “hardcoded” into process binaries

How to avoid collisions across different processes?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation

STATIC: LAYOUT IN MEMORY



STATIC RELOCATION: DISADVANTAGES

No protection

- Process can destroy OS or other processes
- No privacy

Cannot move address space after it has been placed

- May not be able to allocate new process

REPEAT: HOW TO VIRTUALIZE MEMORY?

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions across different processes?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Dynamic Relocation: Base
4. Base+Bounds
5. Segmentation

3) DYNAMIC RELOCATION


Goal: Protect processes from one another

Requires hardware support

- Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates logical or virtual addresses (in their address space)
- Memory hardware uses physical or real addresses



HARDWARE SUPPORT FOR DYNAMIC RELOCATION

Two operating modes

- Privileged (protected, kernel) mode: OS runs
 - When enter OS (trap, system calls, interrupts, exceptions)
 - Allows certain instructions to be executed
 - **Can manipulate contents of MMU**
 - **Allows OS to access all of physical memory**
- User mode: User processes run
 - **Perform translation of logical address to physical address**


Minimal MMU contains **base register** for translation


- base: start location for address space

IMPLEMENTATION OF DYNAMIC RELOCATION: BASE REG

Translation on every memory access of user process



- MMU adds base register to logical address to form physical address
- Each process has different value in base register











VISUAL EXAMPLE OF DYNAMIC RELOCATION: BASE REGISTER











	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1
2 KB		load 4196, R1
3 KB		load 5196, R1
4 KB	P2	load 2024, R1
5 KB		
6 KB		


QUIZ: WHO CONTROLS THE BASE REGISTER?

What entity performs translation of addresses using base register?


- (1) process, (2) OS, or (3) HW

What entity controls the contents and modifies the base register?

- (1) process, (2) OS, or (3) HW



How well does dynamic relocation do with base register for protection?



How well does dynamic relocation do with base register for protection?

BREAK!

Ask your neighbors...

How is Project 1 coming along? What do they think about the discussion videos?
Have they run any of the tests for P1a yet? What seems to be the hardest aspect?

4) DYNAMIC WITH BASE+BOUNDS

Idea: limit the address space with a bounds register

Base register: smallest physical addr (or starting location)

Bounds register: size of this process's virtual address space



- Sometimes defined as largest physical address (base + size)



OS kills process if process loads/stores beyond bounds

IMPLEMENTATION OF BASE+BOUNDS

Translation on every memory access of user process

- MMU compares logical address to bounds register
 - if logical address is greater, then generate error
- MMU adds base register to logical address to form physical address





MANAGING PROCESSES WITH BASE AND BOUNDS

Context-switch

- Add base and bounds registers to PCB
- Steps
 - Change to privileged mode
 - Save base and bounds registers of old process
 - Load base and bounds registers of new process
 - Change to user mode and jump to new process

What if don't change base and bounds registers when switch? Threads!


Protection requirement

- User process cannot change base and bounds registers
- User process cannot change to privileged mode

BASE AND BOUNDS ADVANTAGES

Advantages


- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
 - Can place process at different locations initially and also move address spaces
- Simple, inexpensive implementation
 - Few registers, little logic in MMU
- Fast
 - Add and compare in parallel



BASE AND BOUNDS DISADVANTAGES

Disadvantages

- Each process must be allocated contiguously in physical memory
 - Must allocate memory that may not be used by process
- No partial sharing: Cannot share limited parts of address space




5) SEGMENTATION

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
 - code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected (separate read/write/execute bits)



SEGMENTED ADDRESSING


Process now specifies segment and offset within segment

How does process designate a particular segment?

- Use part of logical address
 - Top bits of logical address select segment
 - Low bits of logical address select offset within segment

What if small address space, not enough bits?

- Implicitly by type of memory reference
- Special registers



SEGMENTATION IMPLEMENTATION

MMU contains Segment Table (for currently running process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:
1 hex digit->4 bits

QUIZ: ADDRESS TRANSLATIONS WITH SEGMENTATION

MMU contains Segment Table (for currently running process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?


Segment	Base	Bounds	R W
0	0x2000	0x6ff	1 0
1	0x0000	0x4ff	1 1
2	0x3000	0xffff	1 1
3	0x0000	0x000	0 0



remember:
1 hex digit->4 bits



Translate logical addresses (in hex) to physical addresses


0x0240: 0x2240
0x1108: 0x0108
0x265c: 0x365c
0x3002: ERROR!

VISUAL INTERPRETATION










ADVANTAGES OF SEGMENTATION

- Enables sparse allocation of address space
 - Stack and heap can grow independently
 - Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc calls sbrk())
 - Stack: OS recognizes reference outside legal segment, extends stack implicitly
- Different protection for different segments
 - Read-only status for code
- Enables sharing of selected segments
- Supports dynamic relocation of each segment



DISADVANTAGES OF SEGMENTATION

Each segment must be allocated contiguously

- May not have sufficient physical memory for large segments

Fix in next lecture with paging...

CONCLUSION

HW+OS work together to virtualize memory

- Give illusion of private address space to each process

Add MMU registers for base+bounds so translation is fast

- OS not involved with every address translation, only on context switch or errors

Dynamic relocation with segments is good building block

- Next lecture: Solve fragmentation with paging